

Vestiaire

Where machine learning meets luxury: decoding the future of second-hand fashion.



Continuing...

Vestiaire Collective is a global platform for buying and selling pre-owned luxury and designer fashion.

- Revenue Streams: Commission-Based Model (once sold, the platform deducts a commission (e.g., 10-40% before paying the seller.)
- Market Capacity: The global secondhand luxury market value stood at US\$34.39 billion in 2023, and is expected to reach US\$60.55 billion by 2029.
- Problems:\$0.84 billion GMV in 2023, significantly lower than the competitor's \$1.83 billion; low conversion rate per listing





Existing analysis

- 3 of the most important features found out to be discount rate, product like count, & buyers fees rate
- LightGBM model with “product sold” as target.
Accuracy 83%
- Causal model suggested top sellers and regular sellers did better with advanced badges signalling their success.



Sentiment Analysis

Product descriptions act as interaction spaces for sellers to market their products through words. Sentiment analysis suggests:

- 60% descriptions express positive sentiments. Example: "Magnificent boiled wool coat. I bought it in the late 2000s..."
- 10% descriptions express negative sentiments. Example: "HARD TO FIND ERMANNO SCERVINO Navy Puffer Jacket with Lace and Fur Accents 44..."
- 30% descriptions are neutral.



Causal Analysis Pt. 2

Does a seller's commissions reflect in the product description sentiment?

Model	ATE	95% CI	Interpretation
S-Learner (LR)	-0.0226	[-0.0226, -0.0226]	A 2.26 percentage-point drop in compound sentiment when you move from low to high price-diff.
S-Learner (XGB)	-0.0310	[-0.4206, 0.3099]	Point estimate is slightly negative, but the CI is so wide that it overlaps zero by a lot—you can't rule out no effect (or even a positive effect).
S-Learner (MLP)	-0.0210	[-0.2370, 0.2001]	Same story as XGB: small negative point estimate, but uncertainty is so large that it's not distinguishable from zero.

Effect:

- All ATEs hover around zero (<1 percentage point).
- Flexible-model CIs are very wide, crossing zero by substantial margins.
- **Conclusion:** There is no convincing causal effect of description sentiment on the likelihood of sale in your data.



Causal Analysis Pt. 2

Does a seller's commissions reflect in the product description sentiment?

Model	ATE	95% CI	Interpretation
S-Learner (LR)	-0.0004	[-0.0004, -0.0004]	A negligible 0.04 percentage-point decrease in sold-rate for "high" sentiment.
S-Learner (XGB)	0.0014	[-0.0331, 0.0368]	Essentially zero, and CI again overlaps zero heavily.
S-Learner (MLP)	0.0077	[-0.0077, 0.0231]	A tiny positive point estimate, but again indistinguishable from zero.

Effect:

- All ATEs hover around zero (<1 percentage point).
- Flexible-model CIs are very wide, crossing zero by substantial margins.
- **Conclusion:** There is no convincing causal effect of description sentiment on the likelihood of sale in your data.



AutoML Setup and Results

Workspace and Data Preparation

Connected to Azure ML workspace: vestiaire_ml_workspace

Loaded 899,281 rows from Azure Blob Storage (cleaned_data.parquet)

Applied feature engineering:

margin_rate, price_to_earning_ratio, price_per_like, seller_activity_ratio

Selected 23 features and sampled 10,000 rows for faster processing

AutoML Configuration and Execution

Task: Classification

Primary metric: AUC_weighted

Submitted experiment: price_elasticity_automl

AutoML automatically:

Performed data quality checks

Detected class imbalance

Trained multiple models (LightGBM, XGBoost, RandomForest, etc.)

Iteration	Pipeline	AUC	Notes
0	MaxAbsScaler LightGBM	0.8242	Good starting point
4	StandardScalerW rapper LightGBM	0.8313	First major improvement
12	StandardScalerW rapper XGBoostClassifier	0.8335	Slight performance boost
24	StandardScalerW rapper LightGBM	0.8372	Best model found
28	MaxAbsScaler LightGBM	0.8315	Consistently strong performance



Value Proposition

Existing Analysis

New analysis

Architecture

Containerization

AutoML Results – Best Model

Metric	Value
Accuracy	0.9847
Weighted AUC	0.8372
Macro AUC	0.8372
Micro AUC	0.9949
Weighted Precision	0.9696
Weighted Recall	0.9847
Weighted F1-Score	0.9771
Log Loss	0.0656
Balanced Accuracy	0.5000
Average Precision (Micro)	0.9946
Average Precision (Macro)	0.5792

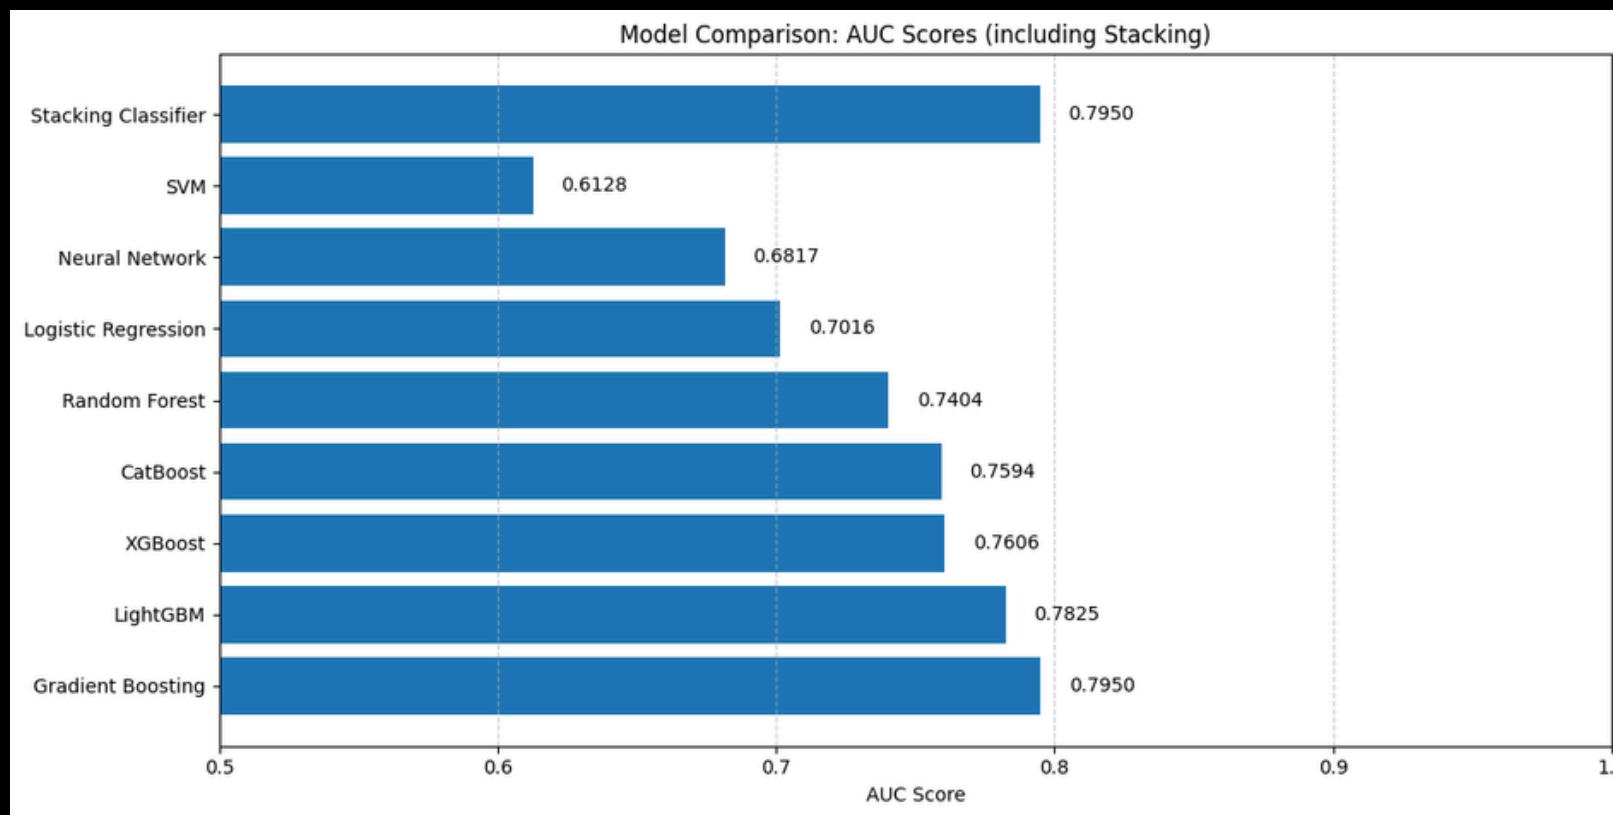
Best Model Performance Summary

- Achieved high overall accuracy of 98.47% and a very strong weighted AUC of 0.8372.
- The model is excellent at predicting the majority class, shown by the high weighted precision (96.96%) and weighted recall (98.47%).
- Micro-AUC score of 0.9949 confirms almost perfect performance across individual samples.
- Low balanced accuracy (0.50) and low macro recall (0.50) indicate the model struggles with the minority class (products that actually sold).
- Despite the class imbalance, the model maintained a very low log loss (0.0656), meaning predictions are highly confident.

The model is extremely reliable for general predictions. What could be done is replicating the ratio of sold to unsold of the whole dataset (instead of 50:50) to emulate real data and have better accuracy.



AutoML Results – Comparison with Baseline Models



Best AUC was achieved by Gradient Boosting and Stacking Classifier (both 0.7950). LightGBM also performed very competitively (AUC 0.7825). Traditional models like Logistic Regression and SVM performed poorly, especially failing to capture the minority class (sold items). Minority class recall was very low across all models (~3-6%), confirming that class imbalance severely affected model sensitivity. Stacking Classifier did not outperform Gradient Boosting individually – ensembling added little in this baseline setup.

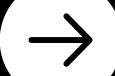
Baseline vs AutoML:

Best baseline AUC ≈ 0.7950

Best AutoML AUC ≈ 0.8372

Model	AUC Score	Recall (Minority Class)	Comment
Gradient Boosting	0.7950	6%	Best standalone model
Stacking Classifier	0.7950	6%	Same as Gradient Boosting
LightGBM	0.7825	6%	Strong, but slightly lower AUC
XGBoost	0.7606	6%	Good performance
CatBoost	0.7594	3%	Good, lower recall
Random Forest	0.7404	3%	Decent baseline
Logistic Regression	0.7016	0%	Failed on minority class
Neural Network	0.6817	3%	Struggled on minority
SVM	0.6128	0%	Weakest performer

AutoML improved AUC by ~5% over manual baselines.



Value Proposition

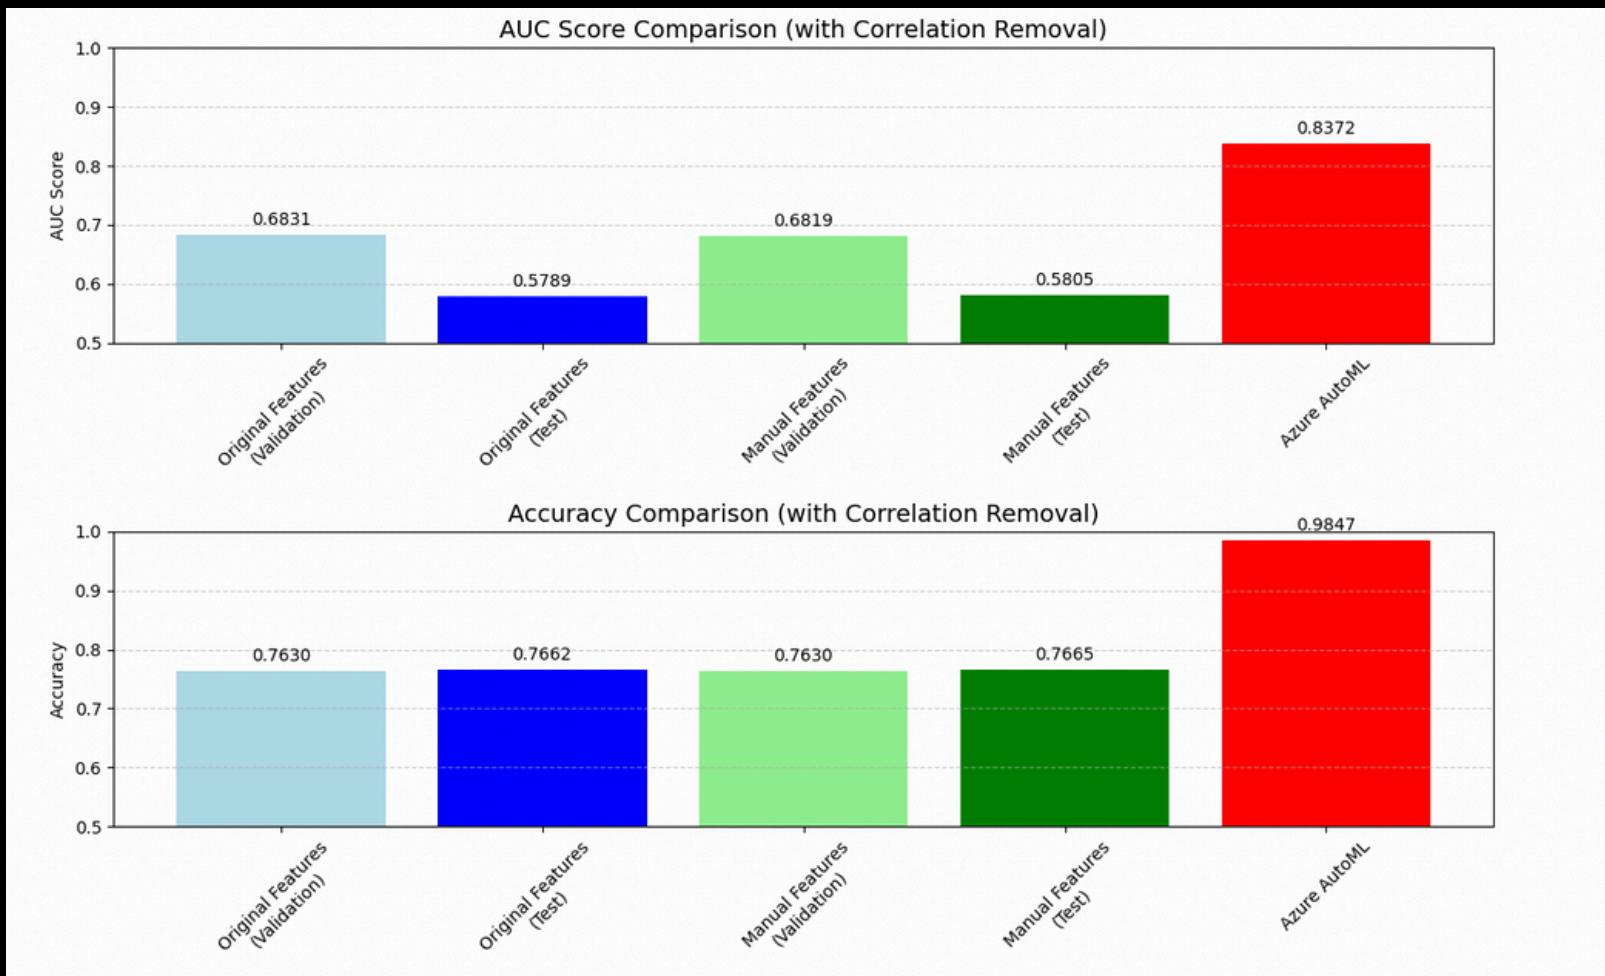
Existing Analysis

New analysis

Architecture

Containerization

AutoML Results – Feature Engineering Comparison



Performed correlation removal (dropped 4 highly correlated features).
Train/Validation/Test split: 60% / 20% / 20%.

Manual feature engineering marginally improved performance over original features (+0.2% AUC).
Both manual and original features struggled to predict minority class despite preprocessing.
AutoML dramatically outperformed manual models (+44% AUC improvement over manual).
Accuracy for AutoML also significantly higher (~98.5% vs ~76%).

Model	AUC	Accuracy (Test)
Original Features (Validation)	0.6831	-
Original Features (Test)	0.5789	0.7662
Manual Features (Validation)	0.6819	-
Manual Features (Test)	0.5805	0.7665
Azure AutoML Best Model	0.8372	0.9847

Manual features alone did not bridge the performance gap.

Azure AutoML's automated feature engineering and modeling pipeline was crucial to achieving top performance.



Value Proposition

Existing Analysis

New analysis

Architecture

Containerization

AutoML Ensemble Experiment

Data Setup:

Data Loaded: 899,281 rows (36 columns)

Sample Used: 15,000 rows (Balanced dataset → 10,000 unsold, 5,000 sold)

After Feature Engineering: 44 features created

After Correlation Removal: 41 features retained

AutoML Configuration:

Task: Classification

Metric: AUC_weighted

Cross-validation: 5 folds

Ensembles: Stacking and Voting enabled after 15 iterations

Max concurrent iterations: 4

Category	Result
Number of child runs	27 models evaluated
Best model type	StackEnsemble
Best model AUC (Validation)	1.0000 (perfect - suspicious)
Best individual model	ExtremeRandomTrees (AUC: ~0.9988)
Ensemble vs Individual Gain	+0.12% AUC improvement
Data Guardrails Check	Passed (Balanced classes, no missing values, no high cardinality)

All models achieved unrealistically perfect AUC scores (close to 1.0), indicating overfitting rather than true generalization.

StackEnsemble and VotingEnsemble showed only marginal gains over individual strong models.

Simple individual models like ExtremeRandomTrees already performed extremely well without needing stacking.

The additional complexity of ensembling appears unnecessary for this dataset.

Root cause may include:
Small and relatively easy-to-separate dataset
Feature engineering capturing too much target signal
Lack of true unseen data variability in sample



Unit Testing

```
import sys
import os
sys.path.insert(0, os.path.abspath(os.path.join(os.path.dirname("__file__"), '../src')))

import seller_analysis_notebook_script as sa
import pandas as pd

def test_no_missing_values():
    df = sa.df
    assert df.isnull().sum().sum() == 0, "There are missing values remaining."

def test_irrelevant_columns_dropped():
    df = sa.df
    dropped_cols = [
        'product_keywords', 'product_id', 'product_name',
        'product_description', 'brand_id', 'brand_url', 'seller_username'
    ]
    for col in dropped_cols:
        assert col not in df.columns, f"Column {col} was not dropped."

print("Unit tests loaded. Run with pytest or call functions manually.")
```

Data Tests

```
import seller_analysis_notebook_script as sa
import pandas as pd

def test_no_duplicate_rows():
    df = sa.df
    assert df.duplicated().sum() == 0, "Duplicate rows found."

def test_column_types_expected():
    df = sa.df
    expected_types = {
        'has_cross_border_fees': 'bool',
        'buyers_fees': 'float64',
    }
    for col, expected_dtype in expected_types.items():
        assert str(df[col].dtype) == expected_dtype, f"{col} is not {expected_dtype}"

print("Data tests loaded. Run with pytest or manually.")
```

Model Validation Tests

```
# Example:
# After training a model, check if the evaluation metrics are acceptable.

def test_model_auc_above_threshold(model_auc):
    assert model_auc > 0.85, "Model AUC is too low."
    print("Model validation test ready (use when model is added).")
```

Testing

Seller Analysis

Load Test

```
import seller_analysis_notebook_script as sa
import pandas as pd

def test_load_big_data():
    df_big = pd.concat([sa.df]*10, ignore_index=True)
    assert df_big.shape[0] == 10 * sa.df.shape[0], "Data not scaled properly."

print("Load test ready.")
```

Data Skew Test

```
import seller_analysis_notebook_script as sa
import pandas as pd

def test_data_skew():
    original_mean = sa.df['buyers_fees'].mean()
    new_data = pd.DataFrame({'buyers_fees': [2.0, 2.5, 3.0]})
    new_mean = new_data['buyers_fees'].mean()
    assert abs(original_mean - new_mean) < 5.0, "Buyers fees distribution shifted significantly."

print("Data skew test ready.")
```

Integration Test

```
import seller_analysis_notebook_script as sa

def test_full_pipeline_runs():
    df = sa.df
    assert df.shape[0] > 0, "DataFrame is empty."
    assert isinstance(sa.num_col, list)
    assert isinstance(sa.cat_col, list)
    assert isinstance(sa.bol_col, list)

print("Integration test ready.")
```

Model Performance Test

```
def test_training_time(model, X_train, y_train):
    start = time.time()
    model.fit(X_train, y_train)
    end = time.time()
    assert end - start < 60, "Model training took too long."
    print("Model performance test ready (use when model is added).")
```

MLflow

Set up MLflow:

```
mlflow.set_tracking_uri("file:///mlruns")
mlflow.set_experiment("Vestiaire_Model_Comparison")
if mlflow.active_run():
    mlflow.end_run()
```

Wrap code

```
with mlflow.start_run(run_name=model_name,
                      nested=True, parent_run_id=parent_run_id):
```

- **Setup:**
 - Just set the tracking URI to tell MLflow where to save logs and artifacts.
- **Wrap Your Code:**
 - Use `mlflow.start_run()` to track everything inside a training run.
- **Log What You Want:**
 - Explicitly log parameters, metrics, and artifacts
- **MLflow handles the rest:**
 - MLflow handles organizing and saving your experiment results automatically.

Specify what to log:

```
mlflow.log_params(model.get_params())
mlflow.log_metrics({
    "roc_auc": roc_auc,
    "accuracy": accuracy,
    "precision": precision,
    "recall": recall,
    "f1_score": f1,
    "log_loss": logloss
})
```

Specify what to log:

```
poetry run mlflow ui
INFO:waitress:Serving on http://127.0.0.1:5000
```

MLflow

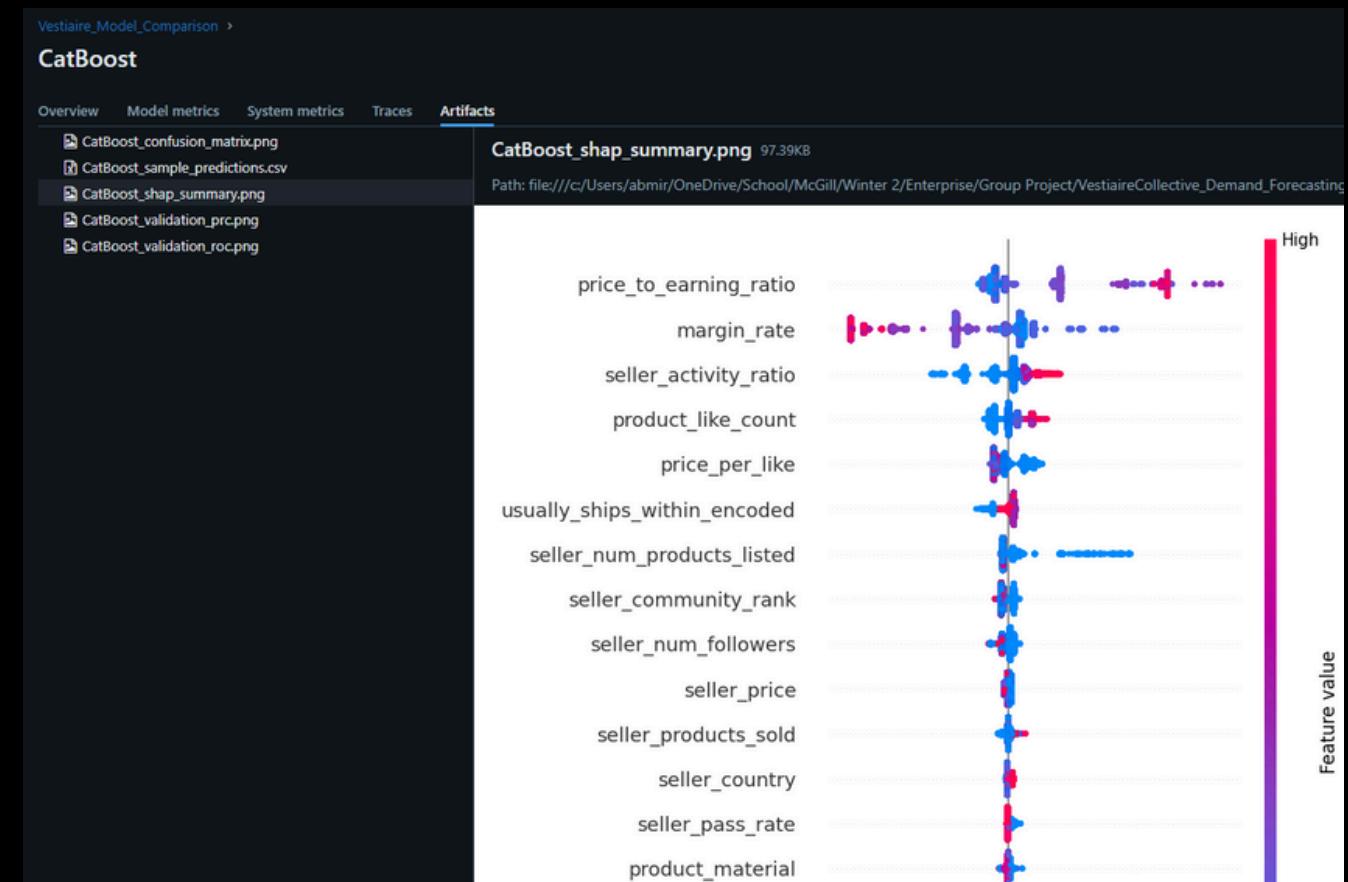
Logs Metrics:

Metrics (18)	
Metric	Value
CatBoost_test_accuracy	0.98760123802287
CatBoost_test_f1	0.3303303303303303
CatBoost_test_log_loss	0.05570689128011065
CatBoost_test_precision	0.9375
CatBoost_test_recall	0.20048602673147023
CatBoost_test_roc_auc	0.8699588756110836
LightGBM_test_accuracy	0.9855811108845933
LightGBM_test_f1	0.13747228381374724
LightGBM_test_log_loss	0.06121900291843374
LightGBM_test_precision	0.7848101265822784
LightGBM_test_recall	0.07533414337788578
LightGBM_test_roc_auc	0.8605019515298695
XGBoost_test_accuracy	0.9870452397279315
XGBoost_test_f1	0.2816032887975334
XGBoost_test_log_loss	0.05760196124114261
XGBoost_test_precision	0.9133333333333333
XGBoost_test_recall	0.16646415552855406
XGBoost_test_roc_auc	0.8679297318442137

Logs Parameters:

Parameters (5)	
Parameter	Value
depth	3
iterations	30
random_state	42
task_type	CPU
verbose	0

Logs Artifacts:



Optuna

- **Bayesian Hyperparameter Tuning with Optuna:**
 - Each Optuna trial is run as a nested MLflow run, sampling hyperparameters (e.g. n_estimators, max_depth, learning_rate) and logging them automatically.
- **Nested MLflow Runs for Trials:**
 - Under a parent run named “Hyperparameter_Tuning,” each trial_i logs its own parameters and validation ROC-AUC metric, making it easy to compare trials in the MLflow UI.
- **Automatic Best-Param Capture:**
 - After all trials, we extract study.best_params and log them (plus the best validation score) in the parent run for transparency and reproducibility.
- **Seamless Transition to Final Training:**
 - We re-initialize our model with the tuned hyperparameters and continue with our usual MLflow-wrapped train/validate/test runs, so everything—tuning and final evaluation—lives in one coherent MLflow experiment.

Run Name
Hyperparameter_Tun...
trial_19
trial_18
trial_17
trial_16
trial_15
trial_14
trial_13
trial_12
trial_11
trial_10
trial_9
trial_8
trial_7
trial_6
trial_5
trial_4
trial_3

From Training to Serving

MLflow Experiment Tracking

- Logged and compared models using MLflow
- Selected best model by ROC AUC and saved with MLflow

MLflow Experiment Tracking

- Deployed the best model in a FastAPI app
- Containerized using Dockerfile_predicting with model config + best_run_id.txt
- Ensures consistent behavior across environments

Local API Testing (Swagger UI)

- Ran live prediction tests from the container.
- Verified model output via Swagger interface. (<http://localhost:8000/docs>)

```
# Load the best model
mlflow.set_tracking_uri("file:/app/mlruns")
# Load best run id
with open("best_run_id.txt", "r") as f:
    best_run_id = f.read().strip()

model = mlflow.pyfunc.load_model(
    model_uri=f"runs:{best_run_id}/model"
)

# FastAPI app
app = FastAPI()
```

```
@app.get("/ping")  ↳ chloe
def ping():
    return {"status": "ok"}

# Input schema
class InputData(BaseModel): 1 usage ↳ chloe
    features: list

❸ Prediction route
@app.post("/predict")  ↳ chloe
def predict(input_data: InputData):
    preds = model.predict([input_data.features])
    return {"prediction": preds.tolist()}
```

The image shows two side-by-side screenshots of the Swagger UI interface for a FastAPI application named 'fastapi-mlflow-app' running on port 8000.

Left Pane (predicting-service):

- Header:** predicting-service, e6ab8528e103, fastapi-mlflow-app, 8000:8000
- Request URL:** http://127.0.0.1:8000/predict
- Server response (Code 200):**

```
{"features": [100.0, 1, 0.85, 10.0, 20, 5.0, 123, 2, 3, 50, 10, 4, 0.3, 5, 0.15, 1, 12, 1, 0, 2, 1, 2, 3]}
```

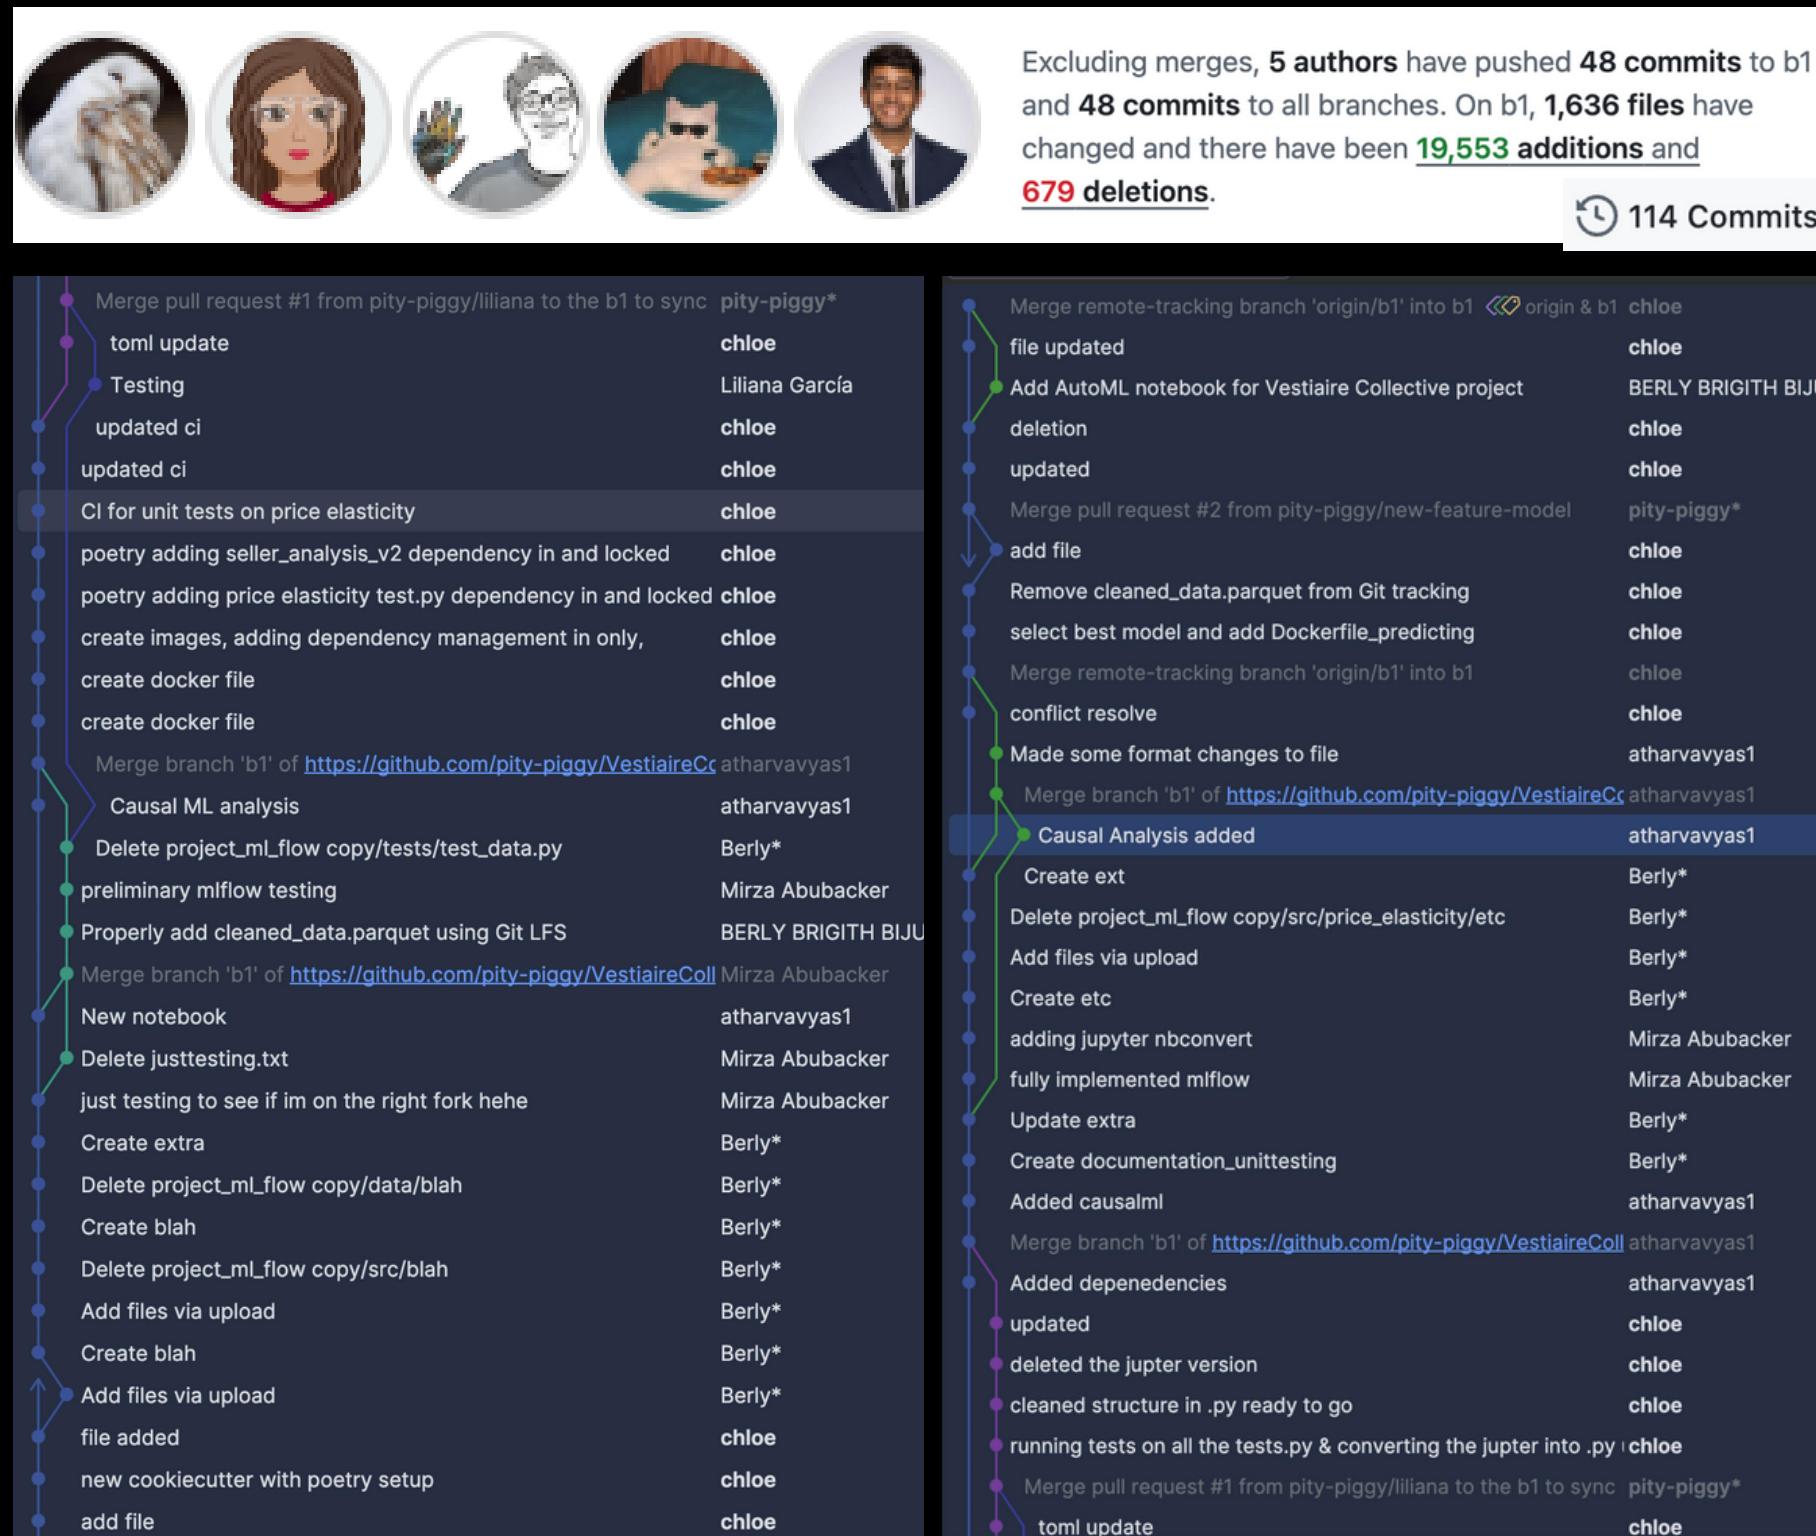
Right Pane (fastapi-mlflow-app):

- Request URL:** http://localhost:8000/predict
- Server response (Code 200):**

```
{"features": [80.0, 1, 0, 0.95, 15.0, 120, 5.5, 10, 3, 2, 25, 50, 0.6, 0.8, 4, 300, 0.3, 1, 20, 1, 2, 3, 1, 5]}
```

Repository Organization & Clarity

We Teamed



We Organized

- **Repo Structure (Cookiecutter)**
 - Ensures scalability, modularity, and ease of navigation for any new developer.
 - **Dependency Management (Poetry)**
 - Environment fully managed with `pyproject.toml` and `poetry.lock` for reproducibility.
 - **Testing with Pytest**
 - All core logic is tested via scripts in the `tests/` folder and integrated into CI for validation.
 - **CI/CD Workflow**
 - Every push or PR to main or b1 triggers automated tests via GitHub Actions — ensuring clean and stable code at every step.
 - **Docker & Containerization**
 - The best-performing model is containerized with a FastAPI app using Docker
- ci_pipeline.yml
- pyproject.toml
- Dockerfile_predicting
- poetry.lock
- ```
auto-ml
data
docs
mlruns
models
module
notebooks
references
reports
src
tests
.gitignore
Dockerfile_predicting
Dockerfile_training
LICENSE
Makefile
README.md
__init__.py
best_run_id.txt
ci_pipeline.yml
poetry.lock
pyproject.toml
```

# Thank You

For Your Attention

# APPENDIX

Value Proposition

Existing Analysis

New analysis

Architecture

Containerization

# Future Considerations



# Model Registration, Deployment, and Monitoring in Azure

McGill University > Vestaire\_ML\_workspace > Endpoints

**Endpoints**

Real-time endpoints Batch endpoints Azure OpenAI Service Serverless endpoints

+ Create Refresh Delete Reset view

| Name                     | Description                         | Quota type | Created on           | Created by              | Updated on           | Compute type       | Compute target | Tags |
|--------------------------|-------------------------------------|------------|----------------------|-------------------------|----------------------|--------------------|----------------|------|
| price-elasticity-service | Price elasticity prediction service |            | Apr 27, 2025 9:04 PM | Berly Berly Bright Blju | Apr 27, 2025 9:04 PM | Container instance |                |      |

McGill University > Vestaire\_ML\_workspace > Endpoints > price-elasticity-service

**price-elasticity-service**

Details Test Consume Logs

**Endpoint attributes**

Service ID: price-elasticity-service  
 Description: Price elasticity prediction service  
 Deployment state: Loading (Running)  
 Operation state: Running  
 Compute type: Container instance  
 Created by: Berly Berly Bright Blju  
 Model ID: price\_elasticity.model3  
 Created on: Apr 27, 2025 9:04 PM  
 Last updated on: Apr 27, 2025 9:04 PM  
 Image ID: ...  
 REST endpoint: /api/v1/predict  
 Key-based authentication enabled: true  
 Swagger URL: ...  
 CPU: 1  
 Memory: 1 GB  
 Application Insights enabled: true  
 Application insights url: https://portal.azure.com/#resource/subscriptions/0ef5d8ea-1c18-4712-8d56-d8f22ed900ea/resourceGroups/VestaireCollective\_Demand\_Forecasting/providers/Microsoft.insights/components/vestairemlwork2885945610

**Tags**: No tags

**Properties**: No data

Command job "plan\_bridge.yaml(2007)" in experiment "prepare\_image" is pending. Job details

McGill University > Vestaire\_ML\_workspace > Jobs

**Jobs**

All experiments All jobs All schedules

Refresh Archive experiment Reset view

| Experiment                      | Last submitted               | Created              | Created by           | Job types                |              |
|---------------------------------|------------------------------|----------------------|----------------------|--------------------------|--------------|
| price_elasticity_automl         | AutoML_85414ea0-6488-4d9c... | Apr 28, 2025 6:59 PM | Apr 27, 2025 9:38 PM | Berly Berly Bright Bl... | Automated ML |
| price_elasticity_mlflow         | mlflow_best_model            | Apr 28, 2025 6:59 PM | Apr 27, 2025 9:38 PM | Berly Berly Bright Bl... | MLflow       |
| autoed_ensemble_comparison      | multi_label_991cpdy          | Apr 28, 2025 3:27 PM | Apr 28, 2025 9:24 PM | Berly Berly Bright Bl... | Automated ML |
| multi_model_ensemble_comparison | Multiclass Classifier        | Apr 28, 2025 3:18 PM | Apr 28, 2025 9:17 PM | Berly Berly Bright Bl... | MLflow       |
| prepare_image                   | plan_bridge.yaml(2007)       | Apr 27, 2025 9:14 PM | Apr 27, 2025 9:14 PM | Berly Berly Bright Bl... | Command      |

McGill University > Vestaire\_ML\_workspace > Models

**Model List**

+ Register Refresh Delete Archive Deploy Compare (preview) Reset view

Search

| Name                   | Version | Type   | Source         | Experiment              | Job (Run ID)                 | Created on           | Tags                          | Properties        |
|------------------------|---------|--------|----------------|-------------------------|------------------------------|----------------------|-------------------------------|-------------------|
| price_elasticity_model | 4       | CUSTOM | This workspace | price_elasticity_automl | AutoML_85414ea0-6488-4d9c... | Apr 28, 2025 6:56 PM | automl : true                 | task : class: ... |
| price_elasticity_model | 3       | CUSTOM | This workspace | price_elasticity_automl | AutoML_85414ea0-6488-4d9c... | Apr 27, 2025 9:02 PM | experiment : price_elastic... | task : class: ... |
| price_elasticity_model | 2       | CUSTOM | This workspace | price_elasticity_automl | AutoML_85414ea0-6488-4d9c... | Apr 27, 2025 9:01 PM | automl : true                 | task : class: ... |
| price_elasticity_model | 1       | CUSTOM | This workspace | price_elasticity_automl | AutoML_85414ea0-6488-4d9c... | Apr 27, 2025 8:59 PM | automl : true                 | task : class: ... |

